

0.1 Coursera

Week 1

Machine Learning algorithms : Supervised learning Unsupervised learning
Reinforcement Learning Recommender systems

Supervised Learning based on right answers Regression problem: predict continuous valued output (price) Classification problem: output value of 0 or 1 To describe the supervised learning problem slightly more formally, our goal is, given a training set, to learn a function $h : X \rightarrow Y$ so that $h(x)$ is a good predictor for the corresponding value of y . For historical reasons, this function h is called a hypothesis. Seen pictorially, the process is therefore like this:

Unsupervised Learning Unsupervised learning allows us to approach problems with little or no idea what our results should look like. We can derive structure from data where we don't necessarily know the effect of the variables.

We can derive this structure by clustering the data based on relationships among the variables in the data.

With unsupervised learning there is no feedback based on the prediction results. Not told what to do -> find some structure in data Clustering Algorithm not give the algorithm the right answers Cocktail party problem: take a data that seems dated together and separate them (two audio sources),

Hypothesis Function

Cost Function Cost Function measures the accuracy of our hypothesis function.

Gradient descent Algorithm Optimize parameters in hypothesis function to minimize the cost function Convex functions has only a global optima, no local optima (good for working with) "Batch" gradient descent: each step of gradient descent uses all the training examples

Learning rate Controls how big the step is in the optimization process

Week 2

Fix a notation throughout the work m -> the number of training examples n -> the number of features j

$x_j^{(i)}$

$i = 1, \dots, m$

$j = 0, \dots, n$ i are the lines and j the columns. Gradient descent for multiple variables

Feature scale mean normalize data to converge faster (do not apply to inserted x_0 that equals 1 for example). Adjust input values as shown in this

formula.

$$x_i := \frac{x_i - \mu_i}{s_i} \quad (1)$$

where μ_i is the average of all values for feature (i) and s_i is the range of values ($\max(x_i) - \min(x_i)$) or the standard deviation of a given feature i .
get all features between -1 and 1 range or 0 and 1

Gradient descent $\mathcal{O}(n^2)$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad (2)$$

works well when n is large

Vectorized implementation of Gradient descent $\mathcal{O}(n^2)$

$$\theta := \theta - \frac{\alpha}{m} X^T (g(X\theta) - Y) \quad (3)$$

Learning rate how to choose: look graph min J vs No. of iterations. It should decrease at every iteration.

Polynomial regression Create new features to fit a polynomial regression to a linear regression. E.g.

$$x_0 = (\text{feature}) \quad (4)$$

$$x_1 = (\text{feature})^2 \quad (5)$$

Analytic solution (faster convergence) The value of $\theta = (X^T X)^{-1} X^T y$ (normal equation) gives the optimal value of θ that minimizes $J(\theta)$ and doesn't need feature scaling. Don't need to iterate nor choose α . Compute the inverse costs a lot $\mathcal{O}(n^3)$. So it is preferably to use if n is small over the iterative method. Noninvertibility: pinv -i pseudo inverse: calculates the value of the inverse even if it is non invertible.

Week 3

Classification problem: take a bunch of data and outputs discrete values
multi class: several discrete outcomes binary class: 2 outputs (0,1)

Logistic regression (classification algorithm): output is always between 0 and 1. sigmoid function or logistic function: $g(z) = \frac{1}{1+e^{-z}}$. Hypothesis: $h(x) = g(\theta^T x)$ maps any real number to the (0,1) interval. $h(x)$ will give the probability that our output is 1.

Decision boundary: line where the hypothesis equals to 0.5. Separates 2 regions, greater than and less than 0.5 to the hypothesis function. The decision boundaries can be linear or not. The decision boundary is a property

of the hypothesis function. The greater the order of the polynomial, the more complex is the shape.

Cost function: find a function that is convex to arrive in a global minimum.

Optimization algorithms:

- Gradient descent
- Conjugate gradient
- BFGS
- L-BFGS

No need to manually pick α in the 3 last algorithms. They are faster than gradient descent. However, they are more complex

Multiclass classification One vs all (one vs rest)

Avoid overfitting: doesn't add too much features, otherwise the training data will not predict well values. The cost function is almost zero to the training example. To fight this problem: reduce the number of features (select which features to keep), regularization (keep all features but reduce the magnitude of parameters theta), Regularization works well when we have a lot of slightly useful features. Add a regularization term in the end of the cost function. A regularization parameter is set to make a trade off between how well the data is fitted and weakening the parameters theta.

In the end we have a regularized linear regression that works much better than the regular one. It even avoids that the normal equation is non-reversible. The calculus is the same, but we add an extra term to the cost function. To calculate the gradient descent, one must actualize the value of theta 0 and of others thetas, since the added sum takes from indexes 1 to n. The value of theta 0 is known as bias.

$$J(\theta) = \frac{1}{2m} \sum_{n=1}^m (h(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{n=1}^n \theta_j^2 \quad (6)$$

Cost Function in logistic regression

$$J(\theta) = \frac{1}{m} \sum_{n=1}^m [-y^{(i)} \log(h(x^{(i)})) - (1 - y^{(i)}) \log(1 - h(x^{(i)}))] \quad (7)$$